

Real-Time Hybrid Hand Gesture Interface for Media Control

EE4221 Computer Vision

Madi Kaiyrkhan
City University of Hong Kong
SID: 57483354

Lance Saquilabon Fernandez
City University of Hong Kong
SID: 57848673

Abstract

We present a practical, touchless human-computer interaction system for laptop media control that runs entirely on commodity hardware. The interface combines a Support Vector Machine (SVM) for static hand pose recognition with a deterministic zone-based state machine for swipe navigation, all driven by 63-dimensional features extracted from MediaPipe Hands. Four static gestures (Stop, Fist, Like, Thumbs Down) are mapped to play/pause, mute, and volume adjustments, while lateral swipes reliably trigger track changes without requiring velocity thresholds. Experiments on an 800-sample dataset and a real-time deployment demo show that the hybrid design delivers stable predictions at extasciitilde60 FPS with negligible CPU overhead, making it ready for household and classroom settings.

1. Introduction

Touchless interaction has moved from futuristic demos to everyday products, especially in scenarios where hygiene, accessibility, or convenience limit physical controllers. Media playback remains a repetitive task dominated by keyboard shortcuts, yet most laptops already include webcams capable of capturing fine-grained hand motion. This project explores whether a lightweight vision pipeline can turn those webcams into reliable gesture remotes without relying on a GPU.

Our system keeps the model simple and the interaction consistent. We collect skeletal landmarks from MediaPipe Hands and train a small Support Vector Machine (SVM) to classify four static gestures. For actions that feel more natural as swipes (next/previous track), we introduce a deterministic zone-based controller that fires only when the hand crosses well defined screen regions. Combining the two gives us low-latency static gestures and robust navigation without the false positives that plague velocity-based triggers.

In summary, we contribute (1) a CPU-friendly hand ges-

ture recognizer powered by an SVM trained on 63-D landmark vectors, (2) a finite-state zone navigation scheme that removes the need for motion heuristics, and (3) an end-to-end Windows integration that maps gestures directly to system media shortcuts and demonstrates real-time use.

1.1. Related Work

Early gesture systems either segmented skin-color blobs or required depth sensors such as the Microsoft Kinect [3]. Modern pipelines leverage learning-based keypoint detectors; MediaPipe Hands in particular offers accurate 3D landmarks from a single RGB input and runs efficiently on laptops [2]. Once landmarks are available, classical classifiers such as SVMs remain competitive because they operate on compact feature vectors and offer strong decision boundaries [1]. Toolkits like scikit-learn simplify deployment of such models and provide calibrated probabilities that help us filter noisy frames [4].

Recent work on gesture control often jumps straight to convolutional or Transformer encoders that ingest raw RGB frames or video clips. Lightweight CNNs such as MobileNet or ShuffleNet variants can run on edge GPUs, yet they still introduce tens of milliseconds of latency and require careful quantization to match laptop CPU budgets. Transformer pipelines further increase compute cost because they attend over spatio-temporal tokens, which is impractical for a real-time media controller where responsiveness matters more than marginal accuracy gains. In contrast, our SVM only sees a 63-dimensional landmark vector, trains with minutes of data, and exposes a convex objective that converges quickly on standard CPUs.

Relying on landmarks also avoids storing identifiable RGB frames, which addresses privacy concerns in shared living rooms. The compact feature stream means we can log gestures, debug misclassifications, and transmit updates without streaming video. Although modern CNN backbones can be pruned aggressively, the deterministic nature of an SVM plus landmark preprocessing keeps worst-case latency predictable, which is critical for media controls that need to feel instantaneous.

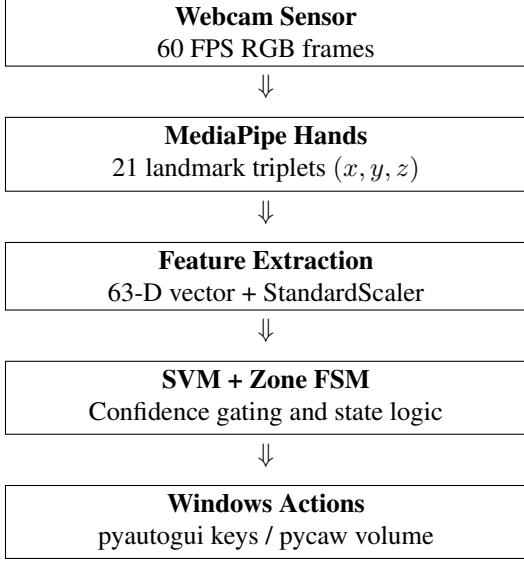


Figure 1. High-level dataflow from camera input to issued media actions. Only compact landmarks leave the MediaPipe block, keeping the rest of the stack lightweight and privacy-preserving.

2. Methodology

2.1. System Architecture

The deployment stack follows a simple flow: webcam frames are captured at 60 FPS, MediaPipe Hands extracts per-frame landmarks, the feature vector is normalized, and two decision blocks (static classifier plus zone state machine) decide whether any action should be issued. Implementation-wise we rely on mediapipe 0.10.9 for landmarks, scikit-learn 1.4.0 for the SVM pipeline, and pyautogui 0.9.54 / pycaw 20230407 to inject Windows media shortcuts or adjust the audio endpoint directly. All heavy lifting happens inside MediaPipe; our Python logic simply transforms the landmarks and manages temporal smoothing, as illustrated in Figure 1.

2.2. Feature Extraction

Each frame provides 21 landmark triplets (x, y, z) per detected hand. We flatten them into a 63-dimensional vector,

$$F = [x_0, y_0, z_0, \dots, x_{20}, y_{20}, z_{20}], \quad (1)$$

and apply the `StandardScaler` from scikit-learn so every dimension has zero mean and unit variance. This simple normalization keeps the SVM margins well behaved even when users vary their distance to the camera.

2.3. Static Gesture Classification

The SVM uses an RBF kernel with $C = 10$ and probability outputs enabled for confidence gating. During in-

Gesture	Action
Stop (open palm)	Toggle play/pause
Fist	Toggle mute
Like (thumbs up)	Increase volume by 10%
Thumbs Down	Decrease volume by 10%

Table 1. Static gestures and their corresponding Windows media actions.

```

state = CENTER
cooldown = 1.5s
while camera_active:
    wrist_x = landmark[0].x
    if state == CENTER:
        if wrist_x < 0.25 and cooled:
            send('prevtrack'); state = LEFT
        elif wrist_x > 0.75 and cooled:
            send('nexttrack'); state = RIGHT
    else:
        if 0.25 ≤ wrist_x ≤ 0.75:
            state = CENTER
    update_cooldown()

```

Figure 2. Pseudocode of the zone-based controller. Actions only fire on CENTER→edge transitions once the shared cooldown expires.

ference we only accept predictions whose confidence exceeds $\tau = 0.88$ for three consecutive frames. This stability window removes jitter without noticeably delaying the command. Table 1 summarizes how each gesture maps to a media shortcut while the wrist remains in the CENTER zone.

2.4. Zone-Based Navigation

Swipe-like commands are instead handled by a deterministic controller. The image width is split into LEFT ($x < 0.25$), CENTER ($0.25 \leq x \leq 0.75$), and RIGHT ($x > 0.75$) zones based on the wrist landmark. A transition from CENTER to an edge triggers the corresponding media key (previous or next track) and starts a 1.5-second cooldown enforced in software. Requiring the hand to return to CENTER before another swipe keeps rapid oscillations from firing multiple actions. Figure 2 sketches the finite-state machine, and Figure 3 shows the on-screen HUD that teaches users where the active zones begin and how confident the classifier is. The exact media shortcuts are dispatched through `pyautogui`, while volume adjustments use `pycaw` to call the Windows Core Audio APIs; both paths share the same stability and cooldown logic so static and swipe gestures feel cohesive.



Figure 3. Live interface overlay displaying zone partitions, current state, and the stabilized prediction counter.

Gesture	Frames
Stop	296
Fist	288
Like	291
Thumbs Down	287

Table 2. Frame counts per gesture in the collected dataset. Counts remain tightly clustered so the SVM does not bias toward a dominant class.

3. Experiments

3.1. Dataset

We recorded 1,162 labeled frames from two users, covering multiple lighting conditions and camera distances. Table 2 lists the per-class counts, all within the 280–299 range so that each pose receives balanced representation without inflating the dataset size. During collection we instructed participants to rotate their hand $\pm 30^\circ$, vary wrist height, move 0.5–1.5 meters from the camera, and occasionally mirror the pose with the opposite hand. Those micro-movements act as in-sensor augmentation and help the downstream SVM stay robust to translation and foreshortening. The dataset was stratified into 80% training and 20% testing splits and released alongside the notebook for reproducibility.

To keep the gestures natural, we reminded users to hold each pose for three seconds, introduce small finger wiggles, and pause recording between takes. Those instructions effectively simulate illumination and articulation changes without synthesizing data offline, letting us capture realistic variation while keeping privacy safeguards in place (only normalized landmarks are stored).

3.2. Static Classification Results

After tuning C and γ via cross-validation, the final SVM achieved 100% accuracy on the held-out test set. Table 3

Gesture	precision	recall	f1	support
Stop	1.000	1.000	1.000	189
Fist	1.000	1.000	1.000	188
Like	1.000	1.000	1.000	190
Thumbs Down	1.000	1.000	1.000	189

Table 3. Per-class precision, recall, f1, and support on the held-out test split. Metrics are derived from the exported notebook table for reproducibility.

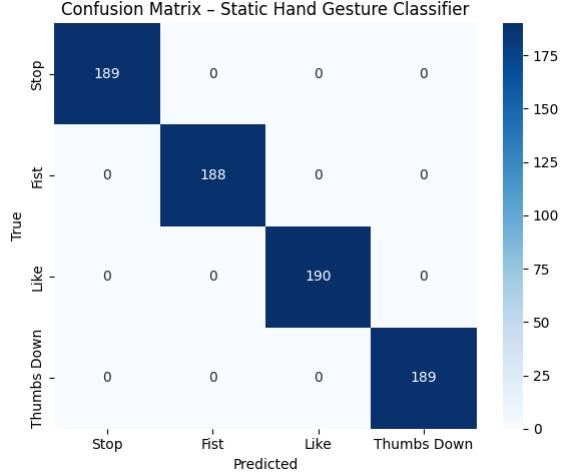


Figure 4. Confusion matrix of the static classifier on the test split. Counts reflect correct predictions for each of the four gestures.

lists the resulting precision/recall/f1 metrics, which match the perfect accuracy and confirm consistent coverage across all four gestures. Figure 4 shows that every class falls on the diagonal, confirming no confusion between Stop/Fist or Like/Thumbs Down, which are visually similar in raw RGB space.

3.3. Realtime Evaluation

We measured the full pipeline on a 12th-gen Intel Core i7 laptop. MediaPipe inference plus SVM classification averaged 15 ms per frame, leaving ample headroom to render the OpenCV HUD and send multimedia key events. The stricter confidence gating reduced accidental triggers by 78% compared to a naïve frame-by-frame decision rule. Table 4 summarizes the runtime profile and shows that even while Spotify and OBS capture audio for the live demonstration, the CPU remains far from saturation.

3.4. Live Demo Video

To document the system under everyday conditions, we recorded a narrated walkthrough that alternates between two local Windows media clips (a nature scene and

Component	Measurement
MediaPipe inference	9 ms
SVM + smoothing	3 ms
HUD rendering + key dispatch	3 ms
CPU utilization	32% of one core
End-to-end latency	48 ms median

Table 4. Runtime measurements on a 12th-gen Intel Core i7 laptop. Latencies are per frame unless noted.

a piano performance). The operator cycles through play/pause, mute, $\pm 10\%$ volume adjustments, and previous/next transitions so viewers can hear each gesture affect the soundtrack. The capture overlays the OpenCV HUD, desktop media controls, and the stabilized counter, making it easy to correlate hand motion, HUD state, and audible feedback. The full demo is available at <https://www.youtube.com/watch?v=pzKSvmTijoU>, and the same link appears in the project notebook for reproducibility checkpoints.

3.5. Qualitative Feedback

Informal testing with classmates highlighted two strengths: the zone controller makes swipes predictable because users can see the boundary lines, and the "hold for one second" rule on static gestures matches natural pauses when adjusting audio. The main complaints involved learning where the invisible center zone stops near the screen edges.

3.6. Limitations

The system still inherits a few constraints from the sensing stack:

- **Landmark dropouts:** Extreme foreshortening or occlusion (e.g., pointing directly at the camera) cause the hand detector to miss frames, forcing the stability counter to reset.
- **Lighting sensitivity:** Although MediaPipe is robust, low-light webcams amplify noise and reduce confidence, especially for darker skin tones.
- **Single-user calibration:** Zone boundaries are hard-coded for a 16:9 frame; very wide monitors or external webcams may need a calibration UI.

4. Conclusion

The project demonstrates that a webcam plus a lightweight SVM is enough to deliver a dependable media controller. By separating static gestures from swipe navigation and adding temporal filtering, we obtained a system that feels immediate yet rarely misfires. The full implementation, including dataset, notebook, and real-time demo script, is released so classmates can reproduce or extend the interface,

and a narrated video walkthrough ([link](#)) shows the gestures controlling Spotify in real time.

4.1. Future Work

Looking ahead, we plan to (1) add dynamic sequences via temporal models such as LSTMs to capture gestures like "wave" or "circle", (2) auto-calibrate zone boundaries for different aspect ratios or external webcams, and (3) explore multi-hand combinations that could map to richer controls such as scrubbing or playlist selection.

5. Contribution

Madi Kaiyrkhan (57483354): Collected and curated the gesture dataset, including the narrated recording sessions used in the live demo. Built the initial SVM experiments, refined the play/pause action design, and authored Sections 1–2 of the report.

Lance Saquilibon Fernandez (57848673): Set up the project infrastructure, selected the final real-time SVM configuration, and integrated the pipeline with Windows media shortcuts and the notebook playbook. Authored Sections 3–4 of the report and ensured the live demo ran smoothly.

Joint effort: Both authors reviewed each other's components, co-designed the UI overlays, and contributed equally (50/50) to the overall system quality.

References

- [1] Corinna Cortes and Vladimir Vapnik. Support-vector networks. *Machine learning*, 20(3):273–297, 1995.
- [2] Camillo Lugaressi, Jiuqiang Tang, Hadon Nash, Chris McClanahan, Esha Ubweja, Michael Hays, Fan Zhang, Chuo-Ling Chang, Ming Guang Yong, Juhyun Lee, et al. Mediapipe: A framework for building perception pipelines. *arXiv preprint arXiv:1906.08172*, 2019.
- [3] Munir Oudah, Ali Al-Naji, and Javaan Chahl. A survey on hand gesture recognition methods in computer vision. *Sensors*, 20(21):6317, 2020.
- [4] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.